

Query Plan Diagnostic Rewrites with Egg

Design of the Egg Plan Transformer Tool

Karl Burtram Student ID: kburtram

June 3, 2022

Introduction

Query execution plans are produced by a database management system to instruct its query processing subsystem how to run a query. A query plan is a directed graph of database operations, such as the join algorithms and access methods to use. The structure of the the execution graph embeds the join order of the tables referenced in the input SQL statement. Optimized query plans are one of the critical aspects determining the overall system performance for a DBMS. Commercial database systems generally produce high-quality plans for common data access patterns. Though creating truly optimal plans is known to be an NP-complete problem, and optimizers often run with tight time constraints, making exhaustive enumeration of all possible plans intractable.

Database administrators and software developers often are asked to investigate why a particular query is taking longer to complete execution than anticipated. A common causes for slow running queries is when the query optimizer generates suboptimal plans that use inefficient processing algorithms given the actual characteristics of the result set being processed. This can occur when there is a significant difference between the amount of data the optimizer expects to be selected based on column statistics and the actual selection cardinalities encountered when the query is executed. This mismatch between estimated and actual cardinalities can lead to undesirable join ordering, join algorithms, and access methods being included in the execution plan. Currently data professionals manually investigate query execution plans to locate these problematic operations and either update the column statistics to better reflect the actual underlying data, or by directly providing query hints to instruct the optimizer to select different operations.

This project introduces a query plan expression cost-based rewrite application, the Egg Plan Transformer (eggplant), that attempts to locate cardinality issues in execution plans, and provide improved query plan expressions. The Egg Plan Transformer is built using the Egg (E-graphs are Good) library that implements equality saturation to discover optimal rewrites. [Willsey, et al., 2020] The optimizations implemented by eggplant are relatively simple, based on a small set of rewrite rules implemented with cardinality-driven cost functions. The rules focus on join algorithm, join ordering, expression tree shape, and access method. Eggplant source code can be reviewed at <https://github.com/kburtram/eggplant>.

The use-case for this project is as part of a query performance tuning workflow to automate the analysis of execution plans. The query plan graph would be translated into a query plan expression and processed by eggplant. Eggplant would apply the rewrite rules with input metadata to find possible improvements on the plan. Commercial database management system query optimizers generate high-quality plans for most SQL statements given accurate statistics. The query rules in eggplant are similar to ones that are already better implemented in SQL Server for the general case, leaving no meaningful optimizations to be further applied. However, since eggplant will have the actual cardinalities it can improve cases where the DBMS choses suboptimal operators based on incorrect column statistics. The rewrites correspond to query hints that can be applied to the the original SQL statement to instruct the DBMS to build a potentially better plan.

This paper briefly introduces common query performance problems, before discussing the design of the Egg Plan Transformer. There is then a background discussion of e-graphs and how eggplant uses cost-based rewrite rules in order to implement optimizations for a subset of the common performance problems.

Common Query Plan Problems

Query performance tuning is one of the most common tasks data professionals are involved in with OLTP systems. [Ben-Gan, 2015] There are numerous sources of query execution inefficiencies that can be detected by analyzing actual query execution plans. Most database systems will allow users to retrieve both estimated and actual query execution plans. Estimated plans show the operations the query optimizer would perform strictly based on statistics. Whereas actual plans show the operations the query optimizer actually performed including actual cardinality results for each operator. Large differences between estimated and actual cardinality estimates are a common cause of suboptimal plans. This is the area where eggplant tries to detect potential plan improvements.

There are many sources of differences between estimated and actual cardinalities. A detailed discussion of these causes is out of scope for this report. For the purposes of building the Egg Plan Transformer component we assume that these bad plans exist and provide citations to additional references for details on why they exist. [Fritchey, 2014] The most common reason for incorrect cardinality estimates is out of date or invalid column statistics. The optimizer uses table statistics to determine join ordering, join algorithm, access methods, and most other query plan operations. Estimates can be out of date in cases such as large tables that are updated frequently, as statistics are typically updated when a certain proportion of the table is changed, which occurs less frequently for large tables. If the new data inserted between automatic statistics updates has a different distribution than the preexisting data than queries for those records can produce inefficient plans. Tables with string, date or geographic indexes can often also have atypical data distributions that cause optimizers to misestimate cardinalities.

Query Execution Plan DSL Representation

Database users typically interact with the query processor with input SQL statements. The DBMS query parser converts this text input into a graphical data structure that models the physical storage engine operations. This execution graph can be usually be retrieved from a DBMS as text output. For the purpose of this project the DBMS query plan representation is manually converted into a simpler DSL representation that is compatible with the Egg rewrite engine. The following BFN language defines this DSL.

```

1 PLAN ::= (select JOIN)
2 JOIN ::= (JOIN_ALGORITHM RELATION_OR_JOIN RELATION_OR_JOIN)
3 RELATION_OR_JOIN ::= RELATION | JOIN
4 RELATION ::= (ACCESS_METHOD IDENTIFIER)
5 JOIN_ALGORITHM ::= hashJoin | mergeJoin | nestedLoopsJoin
6 ACCESS_METHOD ::= scan | seek
7 IDENTIFIER ::= [a-z|A-Z][a-z|A-Z|0-9]
```

Some example query plan expressions and corresponding SQL are provided below.

```

1 # SELECT a.id, b.id
2 # FROM a INNER HASH JOIN b ON a.id = b.id
3 (select (hashJoin (scan a) (scan b)))
4
5 # SELECT a.id, b.id, c.id
6 # FROM a INNER MERGE JOIN b WITH (FORCESEEK) ON a.id = b.id
7 # INNER LOOP JOIN c WITH (FORCESEEK) ON a.id = c.id
8 (select (nestedJoin (mergeJoin (scan a) (seek b)) (seek c)))
```

SQL Statement, Query Plan, and Plan Expression Relationships

The Egg Plan Transformer (eggplant) works with query plan expressions as simplified representations of physical query execution plans which directly represent SQL statements. Data professionals usually prefer to work with SQL-level concepts, such as query plan hints or plan guides, as the starting input and final output of a query tuning workflow. The following example makes the relationship between SQL statements, query execution plans and query plan expressions explicit. And it demonstrates how expression rewrites can be interpreted as query plan hints that can be added to SQL statements to produce the desired query execution plan modifications. [Korotkevitch, 2012]

Consider a simple database schema for the following example with three tables `tbl1`, `tbl2`, `tbl3` that are related to each other through the primary key `id` column in `tbl1`. Assume we join these three tables using the following SQL statement.

```
1 SELECT tbl1.id , tbl2.fid , tbl3.fid
2 FROM tbl1 INNER JOIN tbl2 ON tbl1.id = tbl2.fid
3 INNER JOIN tbl3 ON tbl1.id = tbl3.fid
```

Assume this query is running slower than expected and as part of the query tuning workflow the DBA produces the following query execution plan. The plan shows that the query processor is using nested loops joins for both join algorithms, and is scanning `tbl1` and seeking `tbl2` and `tbl3`.

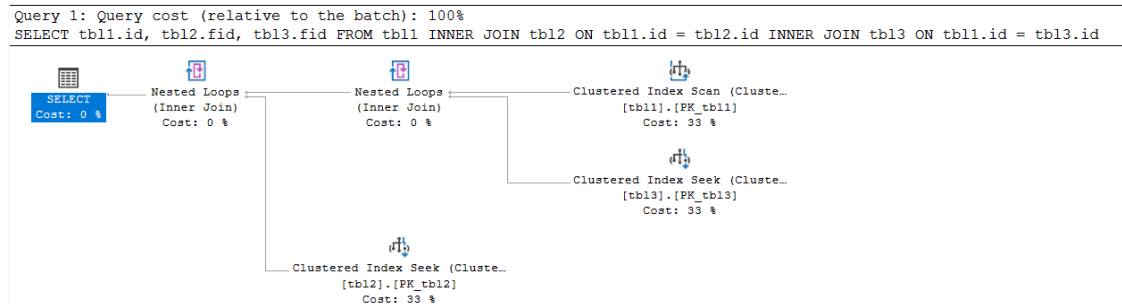


Figure 1: Execution plan for sample query with no join hints

The query plan can be used to create an equivalent eggplant query plan expression and the required input table metadata. In a full end-to-end workflow this process would be automated. The following expression corresponds to the proceeding execution plan.

```
1 (select (nestedLoopsJoin (nestedLoopsJoin (scan tbl1) (seek tbl3)) seek(tbl2)))
```

This expression and the associated metadata can be provided to eggplant to detect inefficiencies. In this simple example, when eggplant processes the rewrite rules using the actual table cardinalities it determines that `tbl1` and `tbl2` should use a hash join and that this intermediate result should be joined with `tbl3` using a merge join. All tables are accessed using a scan. The subsequent sections will provide details on how these rewrite rules and cost functions are implemented.

```
1 (select (mergeJoin (hashJoin (scan tbl1) scan(tbl2)) scan(tbl3)))
```

This query plan expression corresponds to the following T-SQL query. Note the addition of query plan hints to force particular join algorithms and access methods.

```
1 SELECT tbl1.id , tbl2.fid , tbl3.fid
2 FROM tbl1 INNER HASH JOIN tbl2 WITH (FORCESCAN) ON tbl1.id = tbl2.fid
3 INNER MERGE JOIN tbl3 WITH (FORCESCAN) ON tbl1.id = tbl3.fid
```

This updated SQL statement with the added query hints will produce the following physical query execution plan. For this example we assume that these new operations are preferable in this situation based on the actual cardinalities of each of the tables.

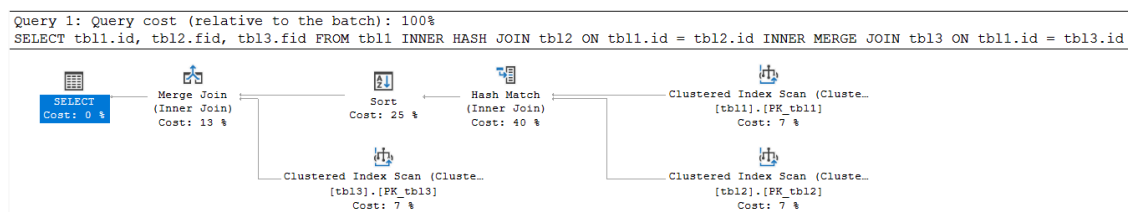


Figure 2: Execution plan for sample query

Egg Plan Transformer Design

The Egg Plan Transformer (eggplant) is implemented as a Rust console application that uses the Egg expression rewriter library. The application accepts a JSON input file that contains a query execution plan expression and metadata about the tables referenced in that expression. Eggplant uses the facilities provided by the Egg library to construct an e-graph and attach the query statistics metadata to e-nodes. There is also a cost function that evaluates various expression rewrite candidates to discover potentially beneficial transformations. The lowest cost equivalent expression is printed as output.

It is important to consider the purpose of the eggplant tool when assessing its design. Eggplant is a component that is intended to be in the inner-core of a query performance tuning workflow, where the input is based on query plans generated from commercial database systems. These database systems have extensive optimization rules, data distribution metrics, and other specialized algorithms. Eggplant is not trying to add new rules that the DBMS is missing. Instead its advantage is that the input metrics contain the actual metrics associated from running the query, not based on estimates from data distribution statistics. This allows eggplant to use the same rules as the DBMS but to produce a different, and potentially better plan, since differences between actual and estimated metrics are a common source of suboptimal query plans. The following diagram provides a high-level illustration of the eggplant inputs and outputs.

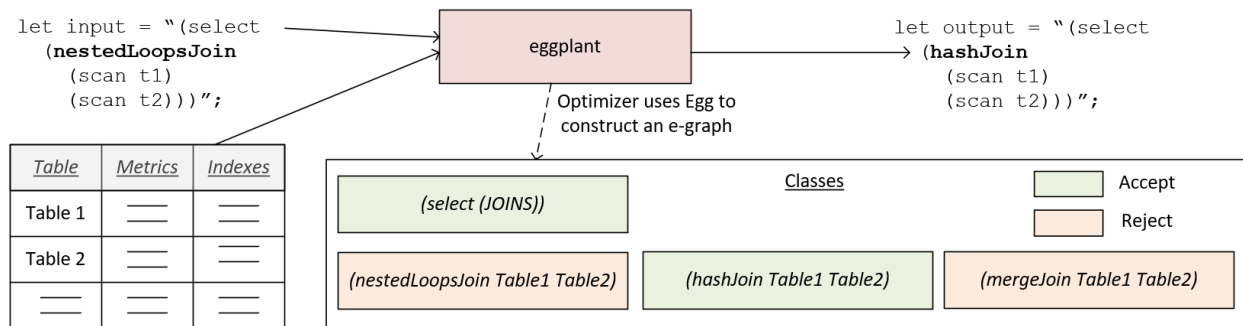


Figure 3: Query execution plan post-processing with Egg-based optimizer

Eggplant accepts an input JSON file that contains the query plan expression and related metadata. Eggplant is a simple expression rewrite tool compared to an actual commercial query optimizer. Therefore there is a limited collection of input metrics, such as actual cardinality, total rows in table, index type, and whether the table is ordered. The following example shows the specific input format.

```
1 {
2   // this is the input query plan expression that eggplant will try to transform
3   "expression": "(select (mergeJoin (scan tbl1) (seek tbl2)))",
4   // metadata about the tables in the expression for use in the cost function
5   "tables": [
6     {
7       "name": "tbl1",           // name of the table in the expression
8       "cardinality": 5,        // the actual cardinality of the table in the query
9       "rows": 1000,            // the number of rows in the table
10      "index": "primary",       // type of index (foreign or primary)
11      "ordered": true           // whether the table will be processed in sorted order
12    },
13    {
14      "name": "tbl2",
15      "cardinality": 45,
16      "rows": 50,
17      "index": "foreign",
18      "ordered": false
19    }
20  ]
21 }
```

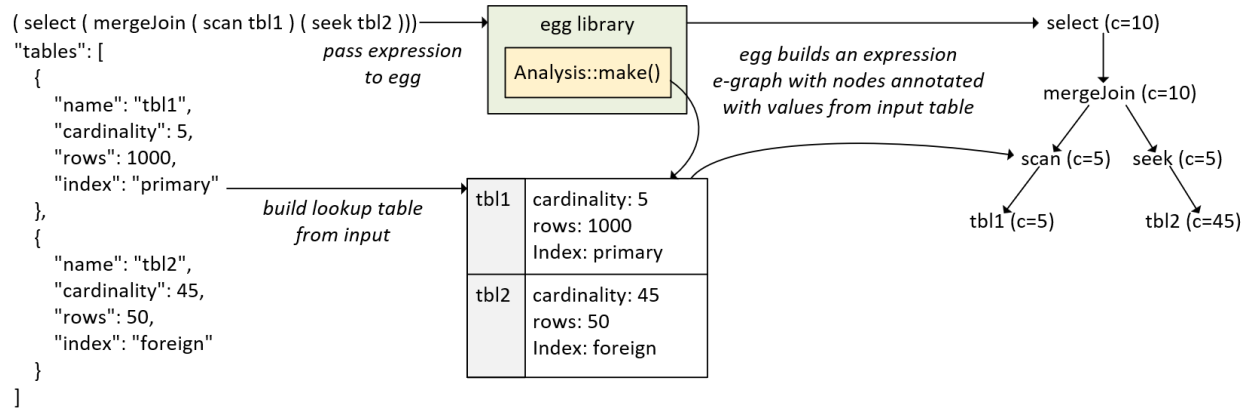


Figure 4: Relationship of input metadata to expression e-graph

The following figure illustrates how an example input expression could be transformed into an improved output expression. First the input statement is transformed into an AST structure. Then e-classes are created for potential rewrite optimizations. In this example the e-classes are for join operator and join order transformations. The join operator is straightforward since the operators are functionally equivalent and can be interchangeable without risk of creating an invalid expression. The join order is more complicated in that the rewrite rules need to ensure that the transformations are still valid. In this specific example `tbl1` can join with either `tbl2` or `tbl3`, but `tbl2` cannot join directly with `tbl3`. After the e-graph rewrites are complete there is a new AST which can be traversed to generate the output expression.

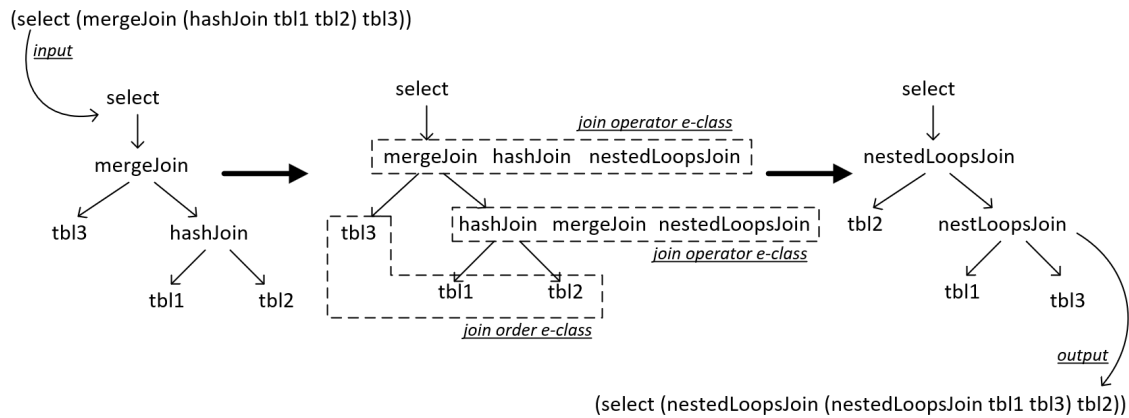


Figure 5: Transformation from input expression to output expression

Eggplant Rules and Cost Functions

Eggplant uses cost-based rewrite rules with the Egg equality saturation engine to discover potential query plan improvements. In each equivalent query plan expression the eggplant cost function will apply context sensitive rules to determine the cost of each node. The cost is generally a combination of the expected subtree cardinality combined with various penalties to prevent undesirable expression rewrites.

The cost function will be given a set of actual cardinalities for each table. It will use these values to provide costs for join operation, join order, and access method. Specifically, whether to join using hash join, merge join, or nested loops join; whether to scan or seek each table; and which order to use to join the tables. There are common guidelines for when to perform each of these optimizations to minimize intermediate result size and limit data I/O to process a query. [West, et al., 2020]

Eggplant supports a limited set of rewrite optimizations as a proof-of-concept to evaluate how well e-graph equality saturation works for cost-based query plan optimizations. The specific set of rules implemented include the following.

- a. Expression tree reordering to create left-deep query plan trees by moving join operators left of table inputs wherever possible. Additionally, tables will be reordered such that tables with primary keys will be moved left of tables with the foreign keys. This transformation simplifies the implementation of the join order rewrite rules.
- b. Join order improvements to reduce the intermediate result set size for plans with multiple join operators. The basic idea is that if we have three tables combined in two joins, such that $t_1 < t_2 < t_3$, then the expression would be changed from $(\text{join}(\text{join } t_1 t_3) t_2)$ to $(\text{join}(\text{join } t_1 t_2) t_3)$.
- c. Access method selection for scanning or seeking the table storage. The criteria is based on the proportion of the actual table cardinality with the number of rows in the table. If $\text{cardinality} / \text{row_count} \geq 0.8$ (i.e. 80% or more of the table) is accessed then a scan is preferred, otherwise, if less than 20% of the table is accessed then a seek is preferred.
- d. Join algorithm improvements to prefer either a hash join, merge join, or nested loops join. This rule is overly simplistic compared to an actual query optimizer but it provides a good illustration of this cost-based rewrite technique. The criteria for each join algorithm is summarized in the below table.

Join Operator	Cost criteria
Hash Join	One small table and one large table. $t_1 \leq 50$ and $t_2 > 1000$
Merge Join	No small tables or tables are ordered. $t_1 > 50$ and $t_2 > 50$ and $t_1 + t_2 > 1000$
Nested Loop Join	Both tables are small. $t_1 + t_2 < 1000$

Determining the desired join order is challenging using the semantics of the Egg cost function methods. Consider the sample of potential arrangements of join operations and tables for the case of two operators and three tables. The position of the nodes in the tree is significant to the Egg rewrite engine. There are 144 distinct arrangements of hash join, merge join, and nested loop join with three tables. Two considerations that arise while using these e-graphs are (1) how to write a minimal set of rewrite rules that will analyze all these expression tree configurations and (2) how to assign unique costs to each of the configurations when we only evaluate a single e-node and its direct descendants at a time in the cost function.

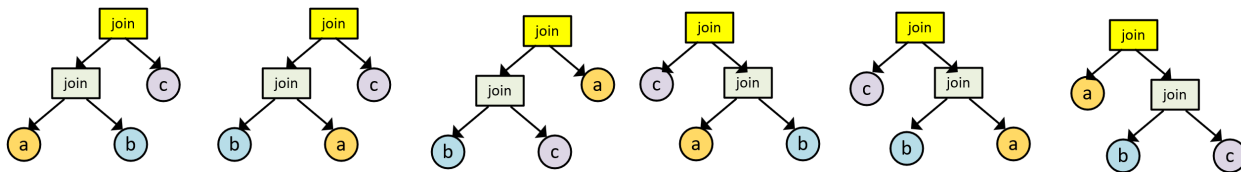


Figure 6: Join Order Table and Algorithm Combinations

The cost functions uses penalties to force preferred expression tree arrangements. The penalties are arbitrary large numbers that can be assigned to nodes in order to exclude that expression from further consideration. If the e-graph can be fully saturated then it is guaranteed that the alternate expression configurations that are not penalized will be visited. The below example illustrates how this process works for determining join order by moving smaller tables lower, and pushing joins and indexes left. Specifically, cardinalities higher in the tree are scaled to 110% cost and out of position primary key (PK) nodes are assigned a 100 cost penalty in this example.

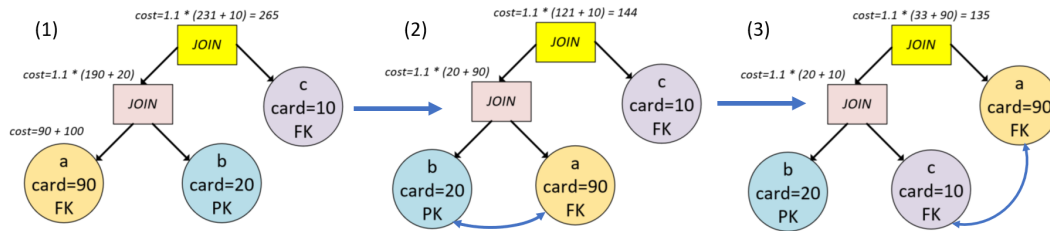


Figure 7: Join Order Costing with Penalties

The following rewrite rules implement transformations for join order, join algorithm, access method, and expression tree shape. The rules only address a small subset of constructs that are possible with ANSI SQL statements. Note that the access method and join ordering rules are only defined for the hashJoin algorithms since Egg will convert other join algorithms to hashJoin during saturation, and then convert back to the preferred join method in the final expression.

Name	Type
idx-left	$(\text{hashJoin } a \ b) \rightarrow (\text{hashJoin } b \ a)$
order-right	$(\text{hashJoin } (\text{hashJoin } a \ b) \ c) \rightarrow (\text{hashJoin } (\text{hashJoin } a \ c) \ b)$
scan-peek	$(\text{hashJoin } (\text{scan } a) \ b) \rightarrow (\text{hashJoin } (\text{seek } a) \ b)$
seek-scan	$(\text{hashJoin } (\text{seek } a) \ b) \rightarrow (\text{hashJoin } (\text{scan } a) \ b)$
hash-join-merge-join	$(\text{hashJoin } a \ b) \rightarrow (\text{mergeJoin } a \ b)$
hash-join-nested-loops-join	$(\text{hashJoin } a \ b) \rightarrow (\text{nestedLoopsJoin } a \ b)$
merge-join-hash-join	$(\text{mergeJoin } a \ b) \rightarrow (\text{hashJoin } a \ b)$
merge-join-nested-loops-join	$(\text{mergeJoin } a \ b) \rightarrow (\text{nestedLoopsJoin } a \ b)$
nested-loops-join-merge-join	$(\text{nestedLoopsJoin } a \ b) \rightarrow (\text{mergeJoin } a \ b)$
nested-loops-join-hash-join	$(\text{nestedLoopsJoin } a \ b) \rightarrow (\text{hashJoin } a \ b)$

The following code snippet is the eggplant hash join cost function implementation. The concept is that the join e-node will be able to access its child e-nodes that contain metrics about the relation, such as cardinality, row count, etc. These metrics will be compared with rules for each operator to assign a penalty to each e-node. The penalty is used as a signal the e-graph extraction algorithm to exclude the expression rewrites that are undesirable.

```

1 // get the costs for the join operations according to join preference rules
2 fn get_hash_join_cost(egraph: &EPlanGraph, table1_id: &Id, table2_id: &Id) -> usize {
3     let mut rank: usize = 800000;
4     let t1_data = &egraph[*table1_id].data;
5     let t2_data = &egraph[*table2_id].data;
6     let t1_cardinality = t1_data.cardinality;
7     let t2_cardinality = t2_data.cardinality;
8     if (t1_cardinality <= 50 || t2_cardinality <= 50)
9         && t1_cardinality + t2_cardinality > 1000 {
10         rank = 2;
11     }
12     if t1_data.index.is_some() {
13         if t1_data.index.as_ref().unwrap() != "primary" {
14             rank += 20000;
15         } else {
16             rank += 10000;
17         }
18     }
19     if t1_data.ordered && t2_data.ordered {
20         rank += 5000;
21     }
22     rank + t1_data.penalty + t2_data.penalty
23 }

```

The complete source code for the eggplant application can be reviewed at <https://github.com/kburtram/eggplant/blob/main/src/main.rs>.

Project Results

This project consists of the background design documented above and the implementation of the eggplant query plan expression optimizer. The evaluation of this system first based on whether the Egg e-graphs are successfully configured to perform the desired rewrites with targeted plan expressions and table metrics. The next evaluation is how well the optimizations reduce the expected size of intermediate results sets from join order and data I/O associated with join operation and access method. The following examples illustrate each of the rewrite rules with sample input metrics.

Join algorithm and access method rewrites Consider the following expression (*select (mergeJoin (scan tbl1) (seek tbl2)))*) that joins two tables that are accessed with scan and seek methods. Given the following three sets of metadata eggplant will rewrite to hashJoin, nestedLoopsJoin and mergeJoin. These examples demonstrate that in these simple cases the rewrites are making the intended transformations based on the cost function rules.

```

1 # Given the below table metrics
2 tbl1 -> cardinality: 5, rows: 1000, index: primary, ordered: true
3 tbl2 -> cardinality: 45, rows: 50, index: foreign, ordered: false
4 # Eggplant will rewrite to use nestedLoopsJoin and alternate scan/seek access methods
5 # Nested loops join is preferable when both tables are small
6 input      (select (mergeJoin (scan tbl1) (seek tbl2)))
7 output     (select (nestedLoopsJoin (seek tbl1) (scan tbl2)))
8
9 # Given the below table metrics
10 tbl1 -> cardinality: 3000, rows: 5000, index: primary, ordered: true
11 tbl2 -> cardinality: 20, rows: 1000, index: foreign, ordered: false
12 # Eggplant will rewrite to use hashJoin and alternate scan/seek access methods
13 # Hash join is preferred when one table is small and one is large
14 input      (select (mergeJoin (scan tbl1) (scan tbl2)))
15 output     (select (hashJoin (scan tbl1) (seek tbl2)))
16
17 # Given the below table metrics
18 tbl1 -> cardinality: 3000, rows: 5000, index: primary, ordered: true
19 tbl2 -> cardinality: 2000, rows: 5000, index: foreign, ordered: true
20 # Eggplant will not find any optimization opportunities
21 # Merge join is preferred when both tables are large and ordered
22 input      (select (mergeJoin (scan tbl1) (scan tbl2)))
23 output     (select (mergeJoin (scan tbl1) (scan tbl2)))

```

Join ordering and parameter ordering rewrites The rewrite rules implement basic join table reordering and parameter reordering. The parameter reordering will move child join operator e-nodes left of child table e-nodes. This has the effect of creating a left-deep join tree. The parameter reordering will also move the table accessed through its primary key left of tables accessed through a foreign key. This is to simplify the implementation of the join ordering rules while maintaining expression equivalency. Specifically, the rules assume when three tables are joined, one will have a primary key and the other two will have foreign keys. In this case the tables with the foreign keys can be exchanged but the primary key table cannot be exchanged with a foreign key table without creating a non-equivalent query plan. This assumption limits the subset of SQL statements that can be processed with these rewrite rules, as there are other valid ways to join three tables, such as using a different key for each of the join operations.

```

1 # Given the below table metrics
2 tbl1 -> cardinality: 3000, rows: 4000, index: primary, ordered: true
3 tbl2 -> cardinality: 2000, rows: 4000, index: foreign, ordered: false
4 tbl3 -> cardinality: 1500, rows: 2000, index: foreign, ordered: false
5
6 # Eggplant will rewrite to move joins left to create a left-deep plan, and swaps
7 # tbl2 and tbl3 in deepest join based on table size to minimize intermediate results size
8 input      (select (nestedLoopsJoin (scan tbl3) (mergeJoin (scan tbl1) (scan tbl2))))
9 output     (select (mergeJoin (mergeJoin (scan tbl1) (scan tbl3)) (scan tbl2)))

```

Bringing it all together These rules can be combined into arbitrarily complex expressions and Egg will apply e-graph saturation to iteratively improve the expression. The following example is a more complete query plan expression that demonstrates how these rules work together to create significantly transformed equivalent expressions that have a lower cost function.


```

1 # Given the below table metrics
2 tbl1 -> cardinality: 3000, rows: 4000, index: primary, ordered: true
3 tbl2 -> cardinality: 2000, rows: 3000, index: foreign, ordered: false
4 tbl3 -> cardinality: 40, rows: 1000, index: foreign, ordered: false
5 tbl4 -> cardinality: 200, rows: 1000, index: foreign, ordered: false
6 tbl5 -> cardinality: 200, rows: 1000, index: foreign, ordered: false
7
8 # Eggplant will rewrite the input statement significantly apply all the rewrite rules,
9 # including join order, join algorithm, expression tree shape, and access method
10 input      (select(
11             nestedLoopsJoin(
12                 nestedLoopsJoin(
13                     mergeJoin (scan tbl1) (scan tbl2))
14                     (scan tbl3))
15                 (hashJoin (scan tbl4) (scan tbl5))))
16 output     (select(
17             mergeJoin(
18                 mergeJoin(
19                     mergeJoin(
20                         hashJoin (scan tbl1) (seek tbl3))
21                         (scan tbl5))
22                     (scan tbl4))
23                 (scan tbl2)))

```

This example demonstrates the flexibility in the eggplant implementation to make broad improvements to query execution plans while only applying rules at a single e-node at a time. This is possible because Egg enumerates through all the permutations of different expression configurations allowing costs and penalties to propagate through the expression tree. This example expression makes over 5000 calls into cost functions as it applies the rewrite rules to generate candidate expressions. The effect of the cost-based rewrite rules can be observed from comparing the input metrics above and the expression tree transformation below. The final expression tree is left-deep, smaller tables are moved down except the primary key table remains at the bottom, hash joins and merge joins are used based on expected table sizes, and access methods are selected based on the cardinality-to-row count ratio. This illustrates that eggplant is performing the expected optimizations consistently even with more complex query plan expressions.

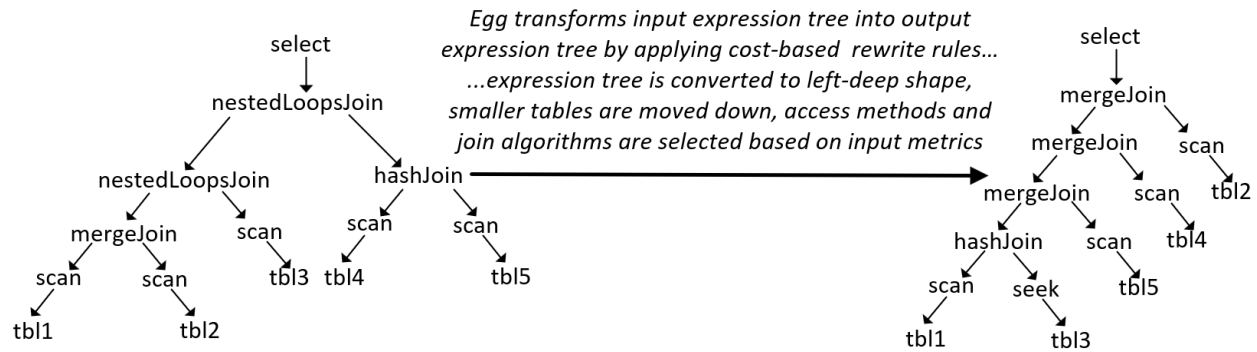


Figure 8: Expression tree transformation for complex query plan expression

The following experiment was used to assess the effectiveness of the join order optimizations. Random expressions and table metrics were generated with table sizes from 2 through 15. These were run through eggplant and the average intermediate result size was collected. Intermediate result size was determined by summing the cardinality of the join operator e-nodes in the e-graph. eggplant was able to find significant improvements in expressions with at least 3 tables, with the improvement leveling off between 40% and 50% after 5 tables. The first chart shows the total intermediate results sizes and the second chart shows the cardinality and penalty improvement percentage. The Python notebook used to run this experiment is available at <https://github.com/kburtram/eggplant/blob/main/analysis/analysis.ipynb>

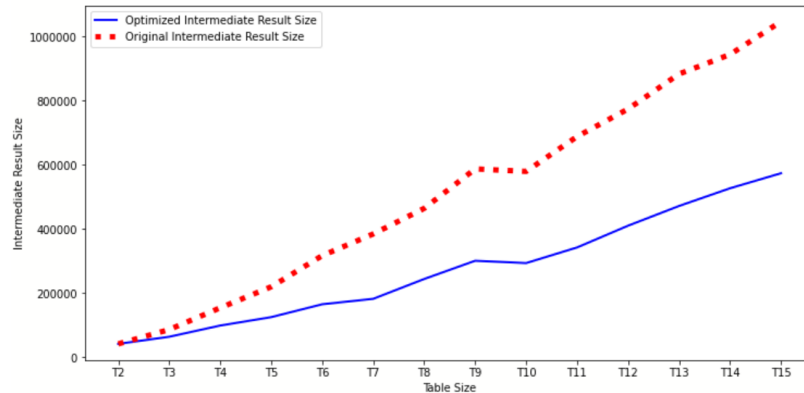


Figure 9: Intermediate Table Results Size by Table Count

Note that as expected the cardinality and penalty improvements are highly correlated since a significant portion of the penalty is computed based on cardinality. The improvements were consistent even as the table count was increased to 50, indicating that with full saturation global optimizations are possible.

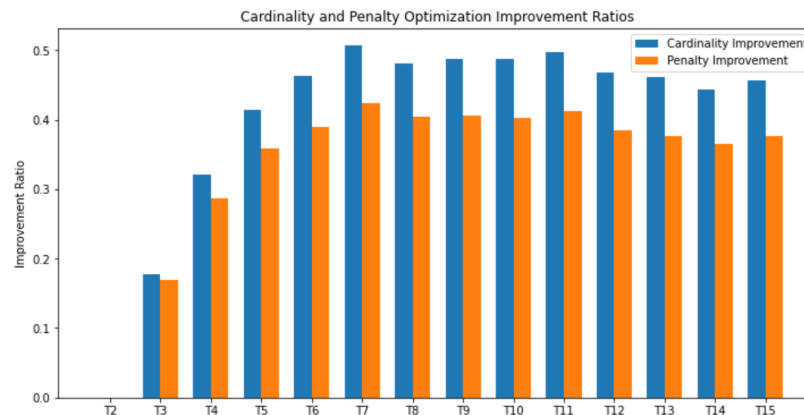


Figure 10: Cardinality and Penalty Improvement by Table Count

Conclusion

The goal of this project was to investigate whether e-graph equality saturation techniques can be used as part of a query performance tuning workflow to identify execution plan inefficiencies. The details around executing SQL queries to generate plans, converting plans to the query plan expression language, and generating query plan hints have been omitted from this project. The project has focused on defining a simple expression language for a subset of physical query plan operators and defining cost-based rewrite rules to implement a small set of query plan optimizations.

The project successfully implemented the eggplant tool that can perform a limited set of query plan expression rewrites based on input table metrics. The project result is neutral in that the eggplant implementation is likely not the ideal mechanism for some of the query plan transformations. Particularly the join ordering rules are relatively limited as only three tables are considered at a time, making global optimizations more difficult. Additionally, the model of using penalties to control the expression rewrites seems to have some limitations when many rules are considered together. The approach generally works since the expression that avoids all the penalties will be selected. The eggplant implementation demonstrates complex parsing and rewriting techniques based on the Egg library. eggplant's general approach could be extended to encompass a more complete set of rules to better improve the query performance tuning workflow.

References

- [Ben-Gan, 2015] Ben Gan, Itzik; Sarka, Dejan; Mechanic, Adam; Farlee, Kevin (2015)
T-SQL Querying
- [Fritchey, 2014] Fritchey, Grant (2014)
SQL Server Query Performance Tuning
- [Korotkevitch, 2012] Korotkevitch, Dmitri (2014)
Pro SQL Server Internals
- [Wang, et al., 2020] Wang, Yisu Remy ; Hutchison, Shana; Leang, Jonathan; Howe, Bill; Suciu, Dan (2020)
SPORES: Sum-Product Optimization via Relational Equality Saturation for Large Scale Linear Algebra
- [West, et al., 2020] West, Randolph; Zecharias, Melody; Assaft, William; Aelterman, Sven; et al. (2020)
SQL Server 2019 Administration
- [Willsey, et al., 2020] Willsey, Max; Nandi, Chandrakana; Wang, Yisu Remy; Flatt, Oliver; et al. (2020)
egg: Fast and Extensible Equality Saturation
- [Yang, et al., 2020] Yang, Yichen; Phothilimthana, Phitchaya Mangpo; Wang, Yisu Remy et al. (2021)
Equality Saturation for Tenor Graph Superoptimization

Source Code and Artifacts

eggplant GitHub - <https://github.com/kburtram/eggplant>

eggplant main source code - <https://github.com/kburtram/eggplant/blob/main/src/main.rs>

eggplant experiment notebook - <https://github.com/kburtram/eggplant/blob/main/analysis/analysis.ipynb>

Presentation slides - <https://github.com/kburtram/eggplant/raw/main/docs/kburtram-presentation.pptx>